drawing

Application of Deep Learning to Text and Image Data

Module 2, Lab 1: Processing Text

In this notebook, you will learn techniques to analyze and process text data. Text processing is known as *natural language processing (NLP)* and is an important topic because of how much information is communicated through text. Knowing how to handle text will help you build models that perform better and are more useful.

You will learn the following:

- What a word cloud is and how to create one
- How to use stemming and lemmatization
- What part-of-speech tagging is and how it impacts text processing
- How to use named entity recognition to sort data

You will be presented with two kinds of exercises throughout the notebook: activities and challenges.

Activity	Challenge
No coding is needed for an activity. You try to understand a concept, answer questions, or run a code cell.	Challenges are where you can practice your coding skills.

Index

Initial Setup

First let's put everything in place.

```
In [1]:   !pip install -U -q -r requirements.txt
```

Install the spaCy library. This will be used to perform some NLP tasks in the lab.

```
In [2]:   !python -m spacy download en_core_web_sm
```

```
/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages/tor
ch/cuda/__init__.py:551: UserWarning: Can't initialize NVML
  warnings.warn("Can't initialize NVML")
Collecting en-core-web-sm==3.7.1
  Downloading https://github.com/explosion/spacy-models/releases/download/e
n_core_web_sm-3.7.1/en_core_web_sm-3.7.1-py3-none-any.whl (12.8 MB)
                                        ━━━━━━━━ 12.8/12.8 MB 99.6 MB/s eta 0:
00:0000:0100:01
Requirement already satisfied: spacy<3.8.0,>=3.7.2 in /home/ec2-user/anacon
da3/envs/pytorch_p310/lib/python3.10/site-packages (from en-core-web-sm==3.
7.1) (3.7.2)
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in /home/ec2-use
r/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.
0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.12)
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in /home/ec2-use
r/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.
0,>=3.7.2->en-core-web-sm==3.7.1) (1.0.5)
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /home/ec2-user/
anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,
>=3.7.2->en-core-web-sm==3.7.1) (1.0.12)
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in /home/ec2-user/anacon
da3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.
2->en-core-web-sm==3.7.1) (2.0.11)
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in /home/ec2-user/anac
onda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.
7.2->en-core-web-sm==3.7.1) (3.0.9)
Requirement already satisfied: thinc<8.3.0,>=8.1.8 in /home/ec2-user/anacon
da3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.
2->en-core-web-sm==3.7.1) (8.2.5)
Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in /home/ec2-user/anaco
nda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.
7.2->en-core-web-sm==3.7.1) (1.1.3)
Requirement already satisfied: srsly<3.0.0,>=2.4.3 in /home/ec2-user/anacon
da3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.
2->en-core-web-sm==3.7.1) (2.5.1)
Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in /home/ec2-user/an
aconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=
3.7.2->en-core-web-sm==3.7.1) (2.0.10)
Requirement already satisfied: weasel<0.4.0,>=0.1.0 in /home/ec2-user/anaco
nda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.
7.2->en-core-web-sm==3.7.1) (0.3.4)
Requirement already satisfied: typer<0.10.0,>=0.3.0 in /home/ec2-user/anaco
nda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.
7.2->en-core-web-sm==3.7.1) (0.9.4)
Requirement already satisfied: smart-open<7.0.0,>=5.2.1 in /home/ec2-user/a
naconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>
=3.7.2->en-core-web-sm==3.7.1) (6.4.0)
Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in /home/ec2-user/anacon
da3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.
2->en-core-web-sm==3.7.1) (4.66.1)
Requirement already satisfied: requests<3.0.0,>=2.13.0 in /home/ec2-user/an
aconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=
3.7.2->en-core-web-sm==3.7.1) (2.31.0)
Requirement already satisfied: pydantic!=1.8,!=1.8.1,<3.0.0,>=1.7.4 in /hom
e/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from s
pacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.10.13)
```

```
Requirement already satisfied: jinja2 in /home/ec2-user/anaconda3/envs/pyto
rch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-core-we
b-sm==3.7.1) (3.1.2)
Requirement already satisfied: setuptools in /home/ec2-user/anaconda3/envs/
pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-cor
e-web-sm==3.7.1) (68.2.2)
Requirement already satisfied: packaging>=20.0 in /home/ec2-user/anaconda3/
envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->e
n-core-web-sm==3.7.1) (21.3)
Requirement already satisfied: langcodes<4.0.0,>=3.2.0 in /home/ec2-user/an
aconda3/envs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=
3.7.2->en-core-web-sm==3.7.1) (3.5.0)
Requirement already satisfied: numpy>=1.19.0 in /home/ec2-user/anaconda3/en
vs/pytorch_p310/lib/python3.10/site-packages (from spacy<3.8.0,>=3.7.2->en-
core-web-sm==3.7.1) (1.23.5)
Requirement already satisfied: language-data>=1.2 in /home/ec2-user/anacond
a3/envs/pytorch_p310/lib/python3.10/site-packages (from langcodes<4.0.0,>=
3.2.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.3.0)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /home/ec2-user/a
naconda3/envs/pytorch_p310/lib/python3.10/site-packages (from packaging>=2
0.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.0.9)
Requirement already satisfied: typing-extensions>=4.2.0 in /home/ec2-user/a
naconda3/envs/pytorch_p310/lib/python3.10/site-packages (from pydantic!=1.
8,!=1.8.1,<3.0.0,>=1.7.4->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (4.8.
0)
Requirement already satisfied: charset-normalizer<4,>=2 in /home/ec2-user/a
naconda3/envs/pytorch_p310/lib/python3.10/site-packages (from requests<3.0.
0,>=2.13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.3.1)
Requirement already satisfied: idna<4,>=2.5 in /home/ec2-user/anaconda3/env
s/pytorch_p310/lib/python3.10/site-packages (from requests<3.0.0,>=2.13.0->
spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /home/ec2-user/anacond
a3/envs/pytorch_p310/lib/python3.10/site-packages (from requests<3.0.0,>=2.
13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.26.18)
Requirement already satisfied: certifi>=2017.4.17 in /home/ec2-user/anacond
a3/envs/pytorch_p310/lib/python3.10/site-packages (from requests<3.0.0,>=2.
13.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (2023.7.22)
Requirement already satisfied: blis<0.8.0,>=0.7.8 in /home/ec2-user/anacond
a3/envs/pytorch_p310/lib/python3.10/site-packages (from thinc<8.3.0,>=8.1.8
->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.7.11)
Requirement already satisfied: confection<1.0.0,>=0.0.1 in /home/ec2-user/a
naconda3/envs/pytorch_p310/lib/python3.10/site-packages (from thinc<8.3.0,>
=8.1.8->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.1.5)
Requirement already satisfied: click<9.0.0,>=7.1.1 in /home/ec2-user/anacon
da3/envs/pytorch_p310/lib/python3.10/site-packages (from typer<0.10.0,>=0.
3.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (8.1.7)
Requirement already satisfied: cloudpathlib<0.17.0,>=0.7.0 in /home/ec2-use
r/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages (from weasel<0.
4.0,>=0.1.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (0.16.0)
Requirement already satisfied: MarkupSafe>=2.0 in /home/ec2-user/anaconda3/
envs/pytorch_p310/lib/python3.10/site-packages (from jinja2->spacy<3.8.0,>=
3.7.2->en-core-web-sm==3.7.1) (2.1.3)
Requirement already satisfied: marisa-trie>=1.1.0 in /home/ec2-user/anacond
a3/envs/pytorch_p310/lib/python3.10/site-packages (from language-data>=1.2-
>langcodes<4.0.0,>=3.2.0->spacy<3.8.0,>=3.7.2->en-core-web-sm==3.7.1) (1.2.
1)
```

✔ Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')

In [3]:
```python
# Import the dependencies
from wordcloud import WordCloud, STOPWORDS
import matplotlib.pyplot as plt
import re, string
import Stemmer
import spacy
from spacy import displacy
import pandas as pd
```

/home/ec2-user/anaconda3/envs/pytorch_p310/lib/python3.10/site-packages/tor
ch/cuda/__init__.py:551: UserWarning: Can't initialize NVML
  warnings.warn("Can't initialize NVML")

Next, you need to create a function to preprocess text so that only real words, not special characters and numbers, are displayed.

In [4]:
```python
# Preprocess text
def preProcessText(text):
    # Lowercase and strip leading and trailing white space
    text = text.lower().strip()

    # Remove HTML tags
    text = re.compile("<.*?>").sub("", text)

    # Remove punctuation
    text = re.compile("[%s]" % re.escape(string.punctuation)).sub(" ", text)

    # Remove extra white space
    text = re.sub("\s+", " ", text)

    # Remove numbers
    text = re.sub(r"[0-9]", "", text)

    return text
```
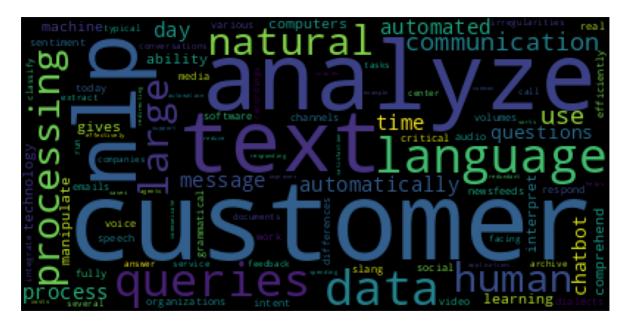
Word cloud

Word clouds, which are also known as *text clouds* or *tag clouds*, help you visualize text data by highlighting the important words or phrases. Word clouds convey crucial information at a glance by making commonly occurring words bigger and bolder. These clouds are commonly used to compare and contrast two pieces of text. Word clouds are also used to identify the topic of a document.

To create a word cloud, you will use WordCloud for Python.

The following text is from the What Is Natural Language Processing (NLP)? page on aws.amazon.com.

In [5]:
```python
text = "Natural language processing (NLP) is a machine learning technology t
ability to interpret, manipulate, and comprehend human language. Organizatic
of voice and text data from various communication channels like emails, text
newsfeeds, video, audio, and more. They use NLP software to automatically pr
the intent or sentiment in the message, and respond in real time to human cc
Natural language processing (NLP) is critical to fully and efficiently analy
It can work through the differences in dialects, slang, and grammatical irre
day-to-day conversations. \
Companies use it for several automated tasks, such as to: \
<li>Process, analyze, and archive large documents</li> \
<li>Analyze customer feedback or call center recordings</li> \
<li>Run chatbots for automated customer service</li> \
<li>Answer who-what-when-where questions</li> \
<li>Classify and extract text</li> \
You can also integrate NLP in customer-facing applications to communicate mc
customers. For example, a chatbot analyzes and sorts customer queries, respc
common questions and redirecting complex queries to customer support. This a
costs, saves agents from spending time on redundant queries, and improves cu
```

In [6]:
```python
# Remove stop words before generating the word cloud
wordcloud = WordCloud(stopwords=STOPWORDS, background_color="black", max_wor

# Clean up the text to prevent plotting punctuation and duplicate words (for
wordcloud.generate(preProcessText(text))

plt.figure(figsize=(15, 10))
plt.axis("off")
plt.imshow(wordcloud);
```

Now that you have created a word cloud, do you see how it can help you quickly identify key words?

Note that the stop words were removed before the graphic was created. This is important so that words that don't impact the meaning of the text aren't overemphasized. Can you think of some examples of stop words?

In this example, you used the precompiled list of stop words that were curated by the WordCloud for Python project. You can print a list of the stop words to make sure that they cover the stop words that you expect.

```
In [7]:  # Show the list of stop words
         ", ".join(list(STOPWORDS))
```

```
Out[7]:  "ourselves, how's, both, would, by, with, com, itself, myself, else, over,
         otherwise, wouldn't, above, doesn't, of, cannot, an, other, whom, you've, b
         een, doing, let's, can't, for, when, himself, you'd, can, in, at, any, was
         n't, after, aren't, because, down, through, your, she'd, while, there's, bu
         t, i'd, such, www, own, i'll, their, are, has, shan't, it, they'd, could, w
         e, he's, only, they're, weren't, when's, theirs, were, isn't, why, them, he
         rself, out, does, do, under, be, me, won't, against, here, also, couldn't,
         some, didn't, the, from, don't, shall, hasn't, that, ours, nor, as, he'd, w
         hat's, mustn't, our, too, hence, she, and, yours, they've, i'm, is, am, no
         t, further, she'll, its, yourselves, haven't, up, so, to, same, until, he,
         more, than, during, we'd, hadn't, each, what, yourself, r, my, a, between,
         very, she's, that's, his, they'll, ought, off, you'll, once, since, which,
         here's, had, where's, most, again, should, where, those, into, who, then, h
         ers, we're, on, you're, this, her, him, these, about, who's, therefore, w
         e've, all, have, being, how, if, it's, no, few, was, why's, he'll, i've, k,
         shouldn't, you, http, or, ever, they, just, before, having, we'll, like, di
         d, below, get, there, themselves, i, however"
```

Part-of-speech tagging

The process of classifying words into their corresponding part of speech based on definition and context is called *part-of-speech tagging*, which is also known as *POS tagging*. A part-of-speech tagger processes a sequence of words and attaches a part-of-speech tag to each word.

For this lab, you will use the the Natural Language Toolkit spaCy. The **nlp() function** creates different token attributes, among them the one representing the token tag: **token.pos**. For example, the following tagged token combines the word *fly* with a noun part of speech tag, *NN*: `tagged_tok = ('fly', 'NOUN')` .

The following table provides the meanings for the tags from a lst of Universal POS tags:

Tag	Meaning
ADJ	Adjective
ADV	Adposition
ADP	Adverb
AUX	Auxiliary
CCONJ	Coordinating Conjuction
DET	Determiner
INTJ	Interjection
NOUN	Noun
NUM	Numeral
PART	Particle
PRON	Pronoun
PROPN	Proper Noun
PUNCT	Punctuation
SCONJ	Subordinating Conjuction
SYM	Symbol
VERB	Verb
X	Other

Now you can use the tagger to tag each token or word in the following text.

Important: Always remember to preprocess the text before tagging, as we have done before in this notebook.

```
In [8]:   # Text sample
          text
```

Out[8]: 'Natural language processing (NLP) is a machine learning technology that gi
        ves computers the ability to interpret, manipulate, and comprehend human la
        nguage. Organizations today have large volumes of voice and text data from
        various communication channels like emails, text messages, social media new
        sfeeds, video, audio, and more. They use NLP software to automatically proc
        ess this data, analyze the intent or sentiment in the message, and respond
        in real time to human communication. Natural language processing (NLP) is c
        ritical to fully and efficiently analyze text and speech data. It can work
        through the differences in dialects, slang, and grammatical irregularities
        typical in day-to-day conversations. Companies use it for several automated
        tasks, such as to: Process, analyze, and archive large documents <
        li>Analyze customer feedback or call center recordings Run chatbot
        s for automated customer service Answer who-what-when-where questi
        ons Classify and extract text You can also integrate NLP in c
        ustomer-facing applications to communicate more effectively with customers.
        For example, a chatbot analyzes and sorts customer queries, responding auto
        matically to common questions and redirecting complex queries to customer s
        upport. This automation helps reduce costs, saves agents from spending time
        on redundant queries, and improves customer satisfaction.'

Try it yourself!

Activity

To use the spaCy part-of-speech tagger, run the following cell.

Observe the tags that are assigned to each word, and use the table from a previous cell

to understand the meaning of each tag.

In [9]:
```python
# Part-of-speech tagging
nlp = spacy.load("en_core_web_sm")
doc = nlp(text)
token_text = [token.orth_ for token in doc]
token_lemma = [token.lemma_ for token in doc]
token_pos = [token.pos_ for token in doc]
pd.DataFrame(zip(token_text, token_pos),
             columns=['token_text', 'token_lemma'])
```

Out[9]:

	token_text	token_lemma
0	Natural	ADJ
1	language	NOUN
2	processing	NOUN
3	(	PUNCT
4	NLP	PROPN
...	...	...
248	and	CCONJ
249	improves	VERB
250	customer	NOUN
251	satisfaction	NOUN
252	.	PUNCT

253 rows × 2 columns

Refer to the table in a previous cell to identify the tags that the spaCy tagger produces.

In [10]:
```
# uncomment the code bellow to display the dependency parse or named entitie
#displacy.render(doc, style="dep")
```

Stemming and lemmatization

Stemming and lemmatization are two ways to process words so that a model will be more efficient. Both methods remove parts of words so that they can be grouped together.

For example, in the following sentence, "ning" would be removed from "running" so that "running" and "run" would be categorized the same.

The child enjoys **running**, so they **run** every day.

What could make stemming and lemmatization difficult to do properly?

Try it yourself!

🖼️Activity

In the next few sections, you will compare stemming and lemmatization.
Consider which text processing method is more suitable for the use case that is provided.

Stemming

Stemming is a rule-based system to convert words into their root forms by removing suffixes. This method helps to enhance similarities (if any) between sentences.

Examples:

"jumping", "jumped" -> "jump"

"cars" -> "car"

```
In [11]:  # let's list the language available
          print(Stemmer.algorithms())
```

```
['arabic', 'armenian', 'basque', 'catalan', 'danish', 'dutch', 'english',
'finnish', 'french', 'german', 'greek', 'hindi', 'hungarian', 'indonesian',
'irish', 'italian', 'lithuanian', 'nepali', 'norwegian', 'porter', 'portugu
ese', 'romanian', 'russian', 'serbian', 'spanish', 'swedish', 'tamil', 'tur
kish', 'yiddish']
```

```
In [12]:  # we will use english for this example
          stemmer = Stemmer.Stemmer('english')
```

```
In [13]:  original_text = "   This is a message to be cleaned. It may involve some thi
          print(original_text)
```

```
    This is a message to be cleaned. It may involve some things like: <br>,
?, :, ''  adjacent spaces and tabs     .
```

```
In [14]:  # Cleaned text
          cleaned_text = preProcessText(original_text)
          print(cleaned_text)
```

```
this is a message to be cleaned it may involve some things like adjacent sp
aces and tabs
```

```
In [15]:  # Use a tokenizer (nlp) and stemmer from the PyStemmer library
          nlp = spacy.load("en_core_web_sm")
          doc = nlp(original_text)
          stemmed_sentence = []
          original_sentence = []
          # Tokenize the sentence
          for token in doc:
              original_sentence.append(token.text)
              stemmed_sentence.append(stemmer.stemWord(token.text))

          pd.DataFrame(zip(original_sentence, stemmed_sentence),
                      columns=['token_text', 'token_stemmer'])
```

Out[15]:

	token_text	token_stemmer
0		
1	This	This
2	is	is
3	a	a
4	message	messag
5	to	to
6	be	be
7	cleaned	clean
8	.	.
9	It	It
10	may	may
11	involve	involv
12	some	some
13	things	thing
14	like	like
15	:	:
16	<	<
17	br	br
18	>	>
19	,	,
20	?	?
21	,	,
22	:	:
23	,	,
24	''	''
25		
26	adjacent	adjac
27	spaces	space
28	and	and
29	tabs	tab
30		
31	.	.
32		

```
In [16]:  stemmed_text = " ".join(stemmed_sentence)
          print(stemmed_text)
```

```
     This is a messag to be clean . It may involv some thing like : < br > ,
? , : , ''   adjac space and tab       .
```

From the output of the previous code cell, you can see that stemming isn't perfect. It makes mistakes, such as "messag", "involv", and "adjac". Stemming is a rule-based method that sometimes mistakenly removes suffixes from words. It does run quickly, which makes it appealing to use for massive datasets.

Lemmatization

If you aren't satisfied with the result of stemming, you can use the lemmatization instead. This method usually requires more work but gives better results.

Lemmatization needs to know the correct word position tags, such as "noun", "verb", or "adjective". You need to use another spaCy function to feed this information to the lemmatizer.

The cell below uses part of the full list of position tags listed in the previous session `Part-of-speech tagging`.

```
In [17]:  # get the original and the lemmatized the word/token
          token_lemma = [token.lemma_ for token in doc]
          token_text = [token.orth_ for token in doc]
          pd.DataFrame(zip(token_text, token_lemma),
                       columns=['token_text', 'token_lemma'])
```

Out[17]:

	token_text	token_lemma
0		
1	This	this
2	is	be
3	a	a
4	message	message
5	to	to
6	be	be
7	cleaned	clean
8	.	.
9	It	it
10	may	may
11	involve	involve
12	some	some
13	things	thing
14	like	like
15	:	:
16	<	<
17	br	br
18	>	>
19	,	,
20	?	?
21	,	,
22	:	:
23	,	,
24	''	''
25		
26	adjacent	adjacent
27	spaces	space
28	and	and
29	tabs	tab
30		
31	.	.
32		

In [18]:
```
lemmatized_text = " ".join(token_lemma)
print(lemmatized_text)
```

```
    this be a message to be clean . it may involve some thing like : < br >
, ? , : , ''   adjacent space and tab       .
```

How do the results compare? Is the lemmatized text better than the stemmed text?

Named entity recognition

Named entity recognition involves identification of key information in text and then classifying that information into predefined categories, such as person, organization, place, or date. This is one of the most popular NLP tasks.

For this section, you will use spaCy. The following table lists the categories and meanings of the category labels that the spaCy module uses.

Category	Meaning
CARDINAL	Numerals that don't fall under another type
DATE	Absolute or relative dates or periods
EVENT	Named hurricanes, battles, wars, sports events, and so on
FAC	Buildings, airports, highways, bridges, and so on
GPE	Countries, cities, states
LANGUAGE	Any named language
LAW	Named documents made into laws
LOC	Non-GPE locations, mountain ranges, bodies of water
MONEY	Monetary values, including unit
NORP	Nationalities, or religious or political groups
ORDINAL	"first", "second", and so on
ORG	Companies, agencies, institutions, and so on
PERCENT	Percentage, including "%"
PERSON	People, including fictional
PRODUCT	Objects, vehicles, foods, and so on (not services)
QUANTITY	Measurements, as of weight or distance
WORK_OF_ART	Titles of books, songs, and so on

The following text was retrieved from the Amazon - The Climate Pledge page of the About Amazon website.

```
In [19]:   # Sample text for named entity recognition
           ner_text = "Amazon and Global Optimism co-founded The Climate Pledge, \
           a commitment to net-zero \
           carbon by 2040."
```

```
In [20]:   # Load the spaCy English pipeline for named entity recognition
           NER = spacy.load("en_core_web_sm")

           # Tag entities in the text
           for word in NER(ner_text).ents:
               print(word.text, word.label_)
```

```
Amazon ORG
Global Optimism ORG
The Climate Pledge WORK_OF_ART
2040 DATE
```

Visualizing the Tags

spaCy has a visualizer called displaCy that you can use to visualize tags. Run the following cell to see it working.

```
In [21]:   # Visual tag with text
           display.render(NER(ner_text), style="ent", jupyter=True)
```

Amazon `ORG` and Global Optimism `ORG` co-founded The Climate Pledge `WORK_OF_ART` , a commitment to net-zero carbon by 2040 `DATE` .

As you can see, named entity recognition can help you identify different entities in text. The process isn't always correct, but it can process large sections of text faster than a human can.

Conclusion

In this lab, you practiced using text processing techniques.

Text visualization aids understanding – Creating a word cloud helps identify the most frequently used words in a dataset, making it a useful tool for quick insights into textual data.

Text normalization improves consistency – Stemming and lemmatization are essential for reducing words to their base forms, improving consistency in text analysis and reducing redundancy in NLP models.

POS tagging enhances text processing – Part-of-speech (POS) tagging provides valuable grammatical context, which is crucial for tasks like sentiment analysis, information retrieval, and machine translation.

NER extracts meaningful information – Named Entity Recognition (NER) enables automatic identification of entities like names, locations, and organizations, making it useful for applications such as customer feedback analysis and document classification.

Preprocessing is key for NLP models – Proper text preprocessing, including tokenization, stemming, and lemmatization, is crucial for building more accurate and efficient machine learning models.

Reducing noise leads to better analysis – Text data often contains noise such as stopwords, special characters, and irrelevant words. Cleaning the data enhances the effectiveness of NLP applications.

Different techniques suit different tasks – While stemming is faster, lemmatization provides more accurate base words. The choice between them depends on the specific NLP application.

Word frequency analysis provides insights – Identifying common words in a dataset can reveal trends, patterns, and potential biases in text data, which is valuable in various fields, from marketing to social sciences.

Linguistic structure impacts interpretation – Understanding the structure of sentences through POS tagging and NER improves text-based decision-making in AI systems, chatbots, and virtual assistants.

NLP techniques have wide applications – The methods covered in this lab are foundational for various real-world applications, such as search engines, recommendation systems, and automated text summarization.

Next lab

In the next lab, you will learn about the bag-of-words (BoW) method to convert text data into numerical values.

End of Lab

In [ ]: